

# Building a Large Scale Dataset for Image Emotion Recognition: The Fine Print and the Benchmark

**Quanzeng You and Jiebo Luo**

Department of Computer Science  
University of Rochester  
Rochester, NY 14623  
[{qyou, jluo}@cs.rochester.edu](mailto:{qyou, jluo}@cs.rochester.edu)

**Hailin Jin**

Adobe Research  
345 Park Avenue  
San Jose, CA 95110  
[hljin@adobe.com](mailto:hljin@adobe.com)

**Jianchao Yang**

Snapchat Inc  
64 Market St  
Venice, CA 90291  
[jianchao.yang@snapchat.com](mailto:jianchao.yang@snapchat.com)

## Abstract

Psychological research results have confirmed that people can have different emotional reactions to different visual stimuli. Several papers have been published on the problem of visual emotion analysis. In particular, attempts have been made to analyze and predict people's emotional reaction towards images. To this end, different kinds of hand-tuned features are proposed. The results reported on several carefully selected and labeled small image data sets have confirmed the promise of such features. While the recent successes of many computer vision related tasks are due to the adoption of Convolutional Neural Networks (CNNs), visual emotion analysis has not achieved the same level of success. This may be primarily due to the unavailability of confidently labeled and relatively large image data sets for visual emotion analysis. In this work, we introduce a new data set, which started from 3+ million weakly labeled images of different emotions and ended up 30 times as large as the current largest publicly available visual emotion data set. We hope that this data set encourages further research on visual emotion analysis. We also perform extensive benchmarking analyses on this large data set using the state of the art methods including CNNs.

## Introduction

Psychological studies have provided evidence that human emotions can be aroused by visual content, e.g. images (Lang 1979; Lang, Bradley, and Cuthbert 1998; Joshi et al. 2011). Based on these findings, recently computer scientists also started to delve into this research topic. However, differently from psychological studies, which mainly focus on studying the changes between physiological and psychological activities of human beings on visual stimuli, most of the works in computer science are trying to predict the aroused human emotion given a particular piece of visual content. Indeed, affective computing, which *aims to recognize, interpret and process human affects* ([http://en.wikipedia.org/wiki/Affective\\_computing](http://en.wikipedia.org/wiki/Affective_computing)), has achieved significant progress in recent years. However,

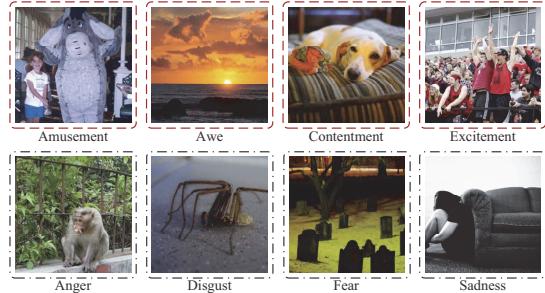


Figure 1: Example images of eight different emotion categories. *Top Row*: four positive emotions and *Bottom Row*: four negative emotions.

the problem of visual emotion prediction is more difficult in that we are trying to predict the emotional reactions given a general visual stimulus, instead of using the collected signals from human's physiological reactions of visual stimuli as studied in affective computing.

The increasing popularity of social networks attracts more and more people to publish multimedia content in online social network platforms. Online users can easily add textual data, e.g., title, descriptions, tags, to their uploaded images and videos. However, these textual information can only help the retrieval of multimedia content in the cognitive level (Machajdik and Hanbury 2010), i.e., semantic level. The meta text data has limited help in bridging the *affective* semantic gap between images pixels and human feelings. In (Hanjalic 2006), the authors call visual emotion prediction *affective content analysis*.

Inspired by the psychology and art theory, different groups of manually crafted features are designed to study the emotional reactions towards visual content. For example, based on art theory (Itten and Van Haagen 1973; Valdez and Mehrabian 1994), Machajdik and Hanbury (Machajdik and Hanbury 2010) defined eight different kinds of pixel level features (e.g. color, texture and composition), which have been empirically proved to be related to emotional reactions. In another recent work (Zhao et al. 2014), *principles-or-art* based features are extracted to classify emotions. Follow-

ing their works, we study the same eight emotions, *Amusement*, *Awe*, *Contentment*, *Excitement*, *Anger*, *Disgust*, *Fear* and *Sadness*. Figure 1 shows the example images for these studied emotions. All these images are selected from the newly constructed data set in this work, where each image is labeled by five Amazon Mechanical Turk workers.

Recently deep learning has enabled robust and accurate feature learning, which in turn produces the state-of-the-art performance on many computer vision related tasks, e.g. digit recognition (LeCun et al. 1989; Hinton, Osindero, and Teh 2006), image classification (Cireşan et al. 2011; Krizhevsky, Sutskever, and Hinton 2012), aesthetics estimation (Lu et al. 2014) and scene recognition (Zhou et al. 2014). One of the main factors that prompt the success of deep learning on these problems is the availability of a large scale data set. From ImageNet (Deng et al. 2009) to AVA dataset (Murray, Marchesotti, and Perronnin 2012) and the very recent Places Database (Zhou et al. 2014), the availability of these data sets have significantly promoted the development of new algorithms on these research areas. In visual emotion analysis, there is no such a large data set with strong labels. More recently, You et al. (You et al. 2015) employed CNNs to address visual sentiment analysis, which tries to bridge the high-level, abstract sentiments concept and the image pixels. They employed the *weakly* label images to train their CNN model. However, they are trying to solve a binary classification problem instead of a multi-class (8 emotions) problem as studied in this work.

In this work, we are interested in investigating the possibility of solving the challenging visual emotion analysis problem. First, we build a large scale emotion data set. On top of the data set, we intend to find out whether or not applying CNNs to visual emotion analysis provides advantages over using a predefined collection of art and psychology theory inspired visual features or visual attributes, which have been done in prior works. To that end, we make the following contributions in this work.

- We collect a large number of weakly labeled emotion related images. Next, we employ Amazon Mechanical Turk to manually label these images to obtain a relatively strongly labeled image data set, which makes the usage of CNN for visual emotion analysis possible. All the data set will be released to the research community upon publishing this work.
- We evaluate the performance of Convolutional Neural Networks on visual emotion analysis and establish it as the baseline for future research. Compared with the state-of-the-art manually crafted visual features, our results suggest that using CNN can achieve significant performance improvement on visual emotion analysis.

## Related Work

Our work is mostly related to both visual emotion analysis and Convolutional Neural Networks (CNNs). Recently, deep learning has achieved massive success on a wide range of artificial intelligence tasks. In particular, deep Convolutional Neural Networks have been widely employed to solve

traditional computer vision related problems. Deep convolutional neural networks typically consist of several convolutional layers and several fully connected layers. Between the convolutional layers, there may also be pooling layers and normalization layers. In early studies, CNNs (LeCun et al. 1998) have been very successful in document recognition, where the inputs are relatively small images. Thanks to the increasing computational power of GPU, it is now possible to train a deep convolutional neural network on large collections of images (e.g. (Krizhevsky, Sutskever, and Hinton 2012)), to solve other computer vision problems, such as scene parsing (Grangier, Bottou, and Collobert 2009), feature learning (LeCun, Kavukcuoglu, and Farabet 2010), visual recognition (Kavukcuoglu et al. 2010) and image classification (Krizhevsky, Sutskever, and Hinton 2012).

However, to the best of our knowledge, there are no related works on using CNNs for visual emotion analysis. Currently, most of the works on visual emotion analysis can be classified into either dimensional approach (Nicolaou, Gunes, and Pantic 2011; Lu et al. 2012) or categorical approach (Machajdik and Hanbury 2010; Borth et al. 2013; Zhao et al. 2014), where the former represents emotion in a continuous two dimensional space and in the later model each emotion is a distinct class. We focus on the categorical approach, which has been studied in several previous works. Jia et al. (2012) extract *color* features from the images. With additional social relationships, they build a factor graph model for the prediction of emotions. Inspired by art and psychology theory, Machajdik and Hanbury (2010) proposed richer hand-tuned features, including *color*, *texture*, *composition* and *content* features. Furthermore, by exploring the principles of art, Zhao et al. (2014) defined more robust and invariant visual features, such as *balance*, *variety*, and *graduation*. Their features achieved the best reported performance on several publicly accessible emotion data sets.

Those hand-tuned visual features have been validated on several publicly available small data sets (see following sections for details). However, we want to verify whether or not deep learning could be applied to this challenging problem and more importantly, on a much larger scale image set. The main issue is that there are no available well labeled data sets for training deep neural networks. Our work intends to provide such a data set for the research community and verify the performance of the widely used deep convolutional neural architecture on this emotion data set.

## Visual Emotion Data Sets

Several small data sets have been used for visual emotion analysis (Yanulevskaya et al. 2008; Machajdik and Hanbury 2010; Zhao et al. 2014), including (1) **IAPS-Subset**: This data set is a subset of the International Affective Picture System (IAPS) (Lang, Bradley, and Cuthbert 1999). This data set is categorized into eight emotional categories as shown in Figure 1 in a study conducted in (Mikels et al. 2005). (2) **ArtPhoto**: Machajdik and Hanbury (2010) built this data set, which contains photos by professional artists. They obtain the ground truth by the labels provided by the owner of each image. (3) **Abstract Paintings**: These are images consisting of both color and

Data Set	Amusement	Anger	Awe	Contentment	Disgust	Excitement	Fear	Sadness	Sum
IAPS-Subset	37	8	54	63	74	55	42	62	395
ArtPhoto	101	77	102	70	70	105	115	166	806
Abstract Paintings	25	3	15	63	18	36	36	32	228
In Total	163	88	171	196	162	196	193	260	1429

Table 1: Statistics of the three existing data sets. The three data sets are imbalanced across the 8 categories.

texture from (Machajdik and Hanbury 2010). They obtain the ground truth of each image by asking people to vote for the emotions of each image in the given eight emotion categories. Table 1 shows the statistics of the existing three data sets. The numbers show that each data set only consists of a very small number of images. Meanwhile, images in all the three different data sets are highly imbalanced. All three data sets are relatively small with images coming from a few specific domains. In particular, for some categories, such as *Anger*, there are less than 10 images. Therefore, if we employ the same methodology (5-fold Cross Validation within each data set) (Machajdik and Hanbury 2010; Zhao et al. 2014), we may have only several images in the training data. This may lead to the possibility that their trained models may have been either over or under fitted.

The above results suggest that the previous efforts on visual emotion analysis deal with small emotion-centric data sets compared with other vision data sets, such as ImageNet (Deng et al. 2009) and Places (Zhou et al. 2014). In this work, we present here a new emotion data set, which is by far the largest available emotion-centric database.

### Building An Image Emotion Dataset from the Wild

There are many different emotion definition systems<sup>1</sup> from psychological and cognitive science. In this work, we use the same eight emotions defined in Table 1, which is derived from a psychological study in (Mikels et al. 2005). Using the similar approach in (Jia et al. 2012), we query the image search engines (Flickr and Instagram) using the eight emotions as keywords. In this way, we are able to collect a total of over 3 million *weakly* labeled images, i.e., labeled by the queries. Next, we delete images which have tags of any two different emotions. We also remove duplicate images using *fdupes*<sup>2</sup>. Figure 2 shows the statistics of the remaining images. As we can see, the number of images in different categories is imbalanced. In particular, there are only small numbers of *contentment* and *disgust* images in both social media platforms. Meanwhile, the number of per category images from Instagram varies significantly. There are much more images from both *Fear* and *Sadness*. This agrees with the finding (<http://goo.gl/vhBBF6>) that people are likely to share sadness from their Instagram accounts.

Next, we employ Amazon Mechanical Turk (AMT) to further label these *weakly* labeled images. In particular, we design a qualification test to filter all workers who want to work on our tasks. The qualification test is designed as an image annotation problem. We randomly select images from

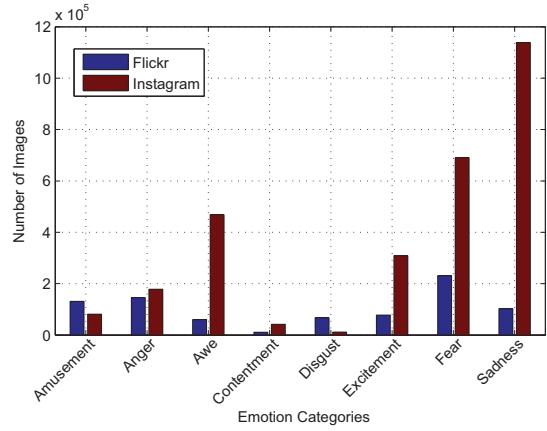


Figure 2: Statistics of the images downloaded from Flickr and Instagram.

the publicly available *ArtPhoto* data set and use the ground-truth labels as the answers. For each given image, we ask the workers to choose the emotion they feel from the eight emotion categories. At first, we conduct experiments within members in our research group. Indeed, the results suggest that this qualification is challenging difficult, in particular when you have to choose only one emotion for each image. Therefore, we design our AMT tasks (HITs) as a much easier verification task instead of the annotation task. Since we have crawled all the images with emotion queries, we have a *weakly* labeled data set. In each HIT, we assign five AMT workers to verify the emotion of each image. For each given image and its *weak* label, we ask them to answer a question like *Do you feel anger seeing this image?* The workers are asked to choose a YES or NO for each image. All workers have to meet the rigorous requirement of correctly answering at least half of the questions (a total of 20) in our qualification test. By the time of finishing work, we have over 1000 workers on our qualification task. Among them, 225 workers meet our qualification criteria.

To start the verification task, we randomly select 11,000 images for each emotion category. After collecting the batch results from AMT, we keep those images which receive at least three Yeses from their assigned five AMT workers. In this way, we are able to build a relatively strongly labeled data set for visual emotion analysis. Table 2 summarizes the number of images for our current data set. The numbers show that different categories may have different acceptance rates for workers to reject or accept the *positive* samples of each verification task. In particular, we add another 2,000

<sup>1</sup><http://en.wikipedia.org/wiki/Emotion>

<sup>2</sup><https://code.google.com/p/fdupes/>

images to make sure that the number of images in *Fear* category is also larger than 1,000 images. In total, we collect about 23,000 images, which is about 30 times as large as the current largest emotion data set, i.e., *ArtPhoto*.

## Fine-tuning Convolutional Neural Network for Visual Emotion Analysis

Convolutional Neural Networks (CNN) have been proven to be effective in image classification tasks, e.g., achieving the state-of-the-art performance in ImageNet Challenge (Krizhevsky, Sutskever, and Hinton 2012). Meanwhile, there are also successful applications by fine-tuning the pre-trained ImageNet model, including recognizing image style (Karayev et al. 2013) and semantic segmentation (Long, Shelhamer, and Darrell 2014). In this work, we employ the same strategy to fine-tune the pre-trained ImageNet reference network (Jia et al. 2014). The same neural network architecture is employed. We only change the last layer of the neural network from 1000 to 8. The remain layers keep the same as the ImageNet reference network. We randomly split the collected 23,000 samples into training (80%), testing (15%) and validating sets (5%).

Meanwhile, we also employ the *weak* labels (Jia et al. 2012) to fine-tune another model as described in (You et al. 2015). We exclude those images that have been chosen to be submitted to AMT for labeling. Next, since *contentment* contains only about 16,000 images, we randomly select 20,000 images for other emotion categories. In this way, we have a total of 156,000 images. We call this model *Noisy-Fine-tuned CNN*. We fine-tune both models using Caffe with a Linux server with 2 NVIDIA TITAN GPUs on top of the pre-trained ImageNet CNN model.

## Performance of Convolutional Neural Networks on Visual Emotion Analysis

After the fine-tuning of the pre-trained CNN model, we obtain two new CNN models. To compare with the ImageNet-CNN, we also show the results of using the SVM trained on features extracted from the second to the last layer of the pre-trained ImageNet-CNN model. In particular, we employ PCA to reduce the dimensionality of the features. We also try several different numbers of principal components. The results are almost the same. To overcome the imbalance problem in the data, we adjust the weights of SVM for different classes (in our implementation, we use LIBSVM<sup>3</sup>, which provides such a mechanism). Table 3 summarizes the performance of the three groups of features on the 15% randomly chosen testing data. The overall accuracy of the Fine-tuned-CNN is almost 60%. As a baseline, the visual features extracted from ImageNet-CNN only lead to an overall accuracy of about 30%, which is half of Fine-tuned-CNN. The Noisy-Fine-tuned-CNN model has an overall accuracy of about 46%, which suggests that this model can learn some knowledge from the noisily labeled images. However, even though it has much more training samples compared with *Fine-tuned CNN*, it fails to outperform *Fine-tuned CNN*, which is trained on strongly labeled samples.

We also calculate the confusion matrix of the three algorithms from their prediction results on the testing data to further analyze their performance. Figure 3(c) shows the confusion matrix of the Fine-tuned CNN model. Compared with the other two models, the true negative rates from *Fine-tuned-CNN* are the best in most emotion categories. Meanwhile, the confusion matrix of *Noisy-Fine-tuned CNN* seems to be more balanced, except for the *contentment* emotion (seeFigure 3(b)). Indeed, these findings are consistent with the number of available labeled samples (see Table 2). The more the labeled images, the higher probability that the corresponding emotion will receive a higher true positive rate. Figure 3(a) shows the confusion matrix using the more general ImageNet-CNN features. It is interesting to see that overall the performance is worse than the Fine-tuned CNN features. However, the true positive rate of *fear* is higher than that of using the Fine-tuned features.

The embedding of the testing images using deep visual features (we do not show the embedding for Noisy-Fine-tuned-CNN due to space arrangement) is shown in Figure 4. The features are also processed using t-SNE (Van der Maaten and Hinton 2008). The embedding using ImageNet-CNN shows that images from the same scene or of similar objects are embedded into neighboring areas. However, the embedding using Figure 4(b) seems to make the images more diverse in terms of objects or scenes. This is indeed comply with the fact that even the same object could lead to different visual emotion at its different state, e.g., angry dog and cute dog.

## Performance of Convolutional Neural Networks on Public Existing Visual Emotion Data Set

We have described several existing data sets in Section . Table 1 summarizes the statistics of the three data sets. To the best of our knowledge, no related studies have been conducted on evaluating the performance of Convolutional Neural Networks on visual emotion analysis. In this section, we evaluate all the three deep neural network models on all the three data sets and compare the results with several other state-of-the-art methods on these data sets.

In particular, we extract deep visual features for all the images in the three data sets using the trained deep neural network models from the second to the last layer. In this way, we obtain a 4096 dimensional feature representation for each image from each deep model. Next, we follow the same evaluation routine described in (Machajdik and Hanbury 2010) and (Zhou et al. 2014). At first, PCA is employed to reduce the dimensions of the features respectively. For all the three data sets, we reduce the number of feature dimensions from 4096 to 20, which is capable of keeping at least 90% variance. Next, a linear SVM is trained on the reduced feature space. Following the same experimental approach, the *one v.s. all* strategy is employed to train the classifier. In particular, we randomly split the data into 5 batches such that 5-fold Cross Validation is used to obtain the results. Also, we assign larger penalties to *true negative* samples in the SVM training stage in order to optimize the *per class true positive rate* as suggested by both (Machajdik and Hanbury 2010) and (Zhou et al. 2014).

<sup>3</sup><http://www.csie.ntu.edu.tw/~cjlin/libsvm/>

Data Set	Amusement	Anger	Awe	Contentment	Disgust	Excitement	Fear	Sadness	Sum
Submitted	11,000	11,000	11,000	11,000	11,000	11,000	13,000	11,000	90,000
Labeled	4,942	1,266	3,151	5,374	1,658	2,963	1,032	2,922	23,308

Table 2: Statistics of the current labeled image data set. Note that all emotion categories have 1000+ images.

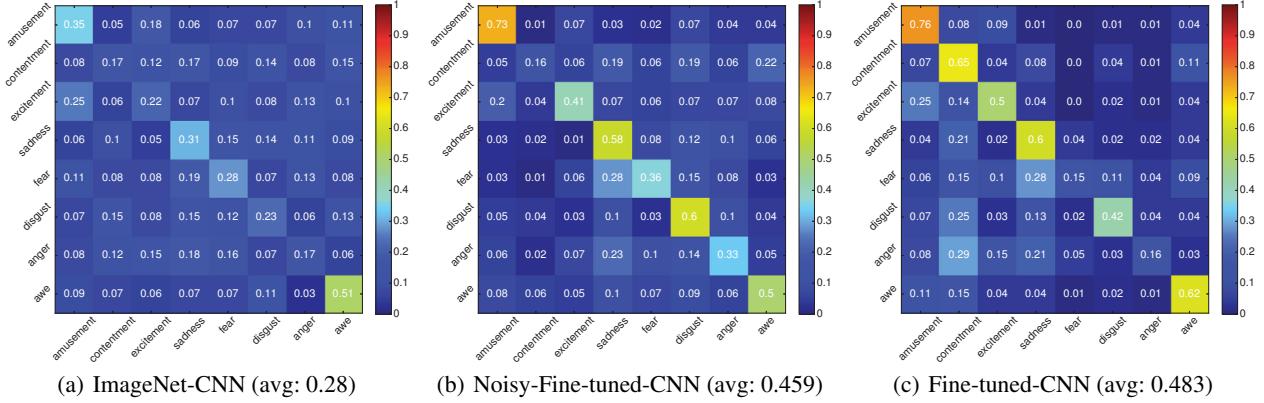


Figure 3: Confusion matrix for ImageNet-CNN, Noisy-Fine-tuned-CNN and Fine-tuned-CNN on the testing Amazon Mechanical Turk (AMT) labeled images.

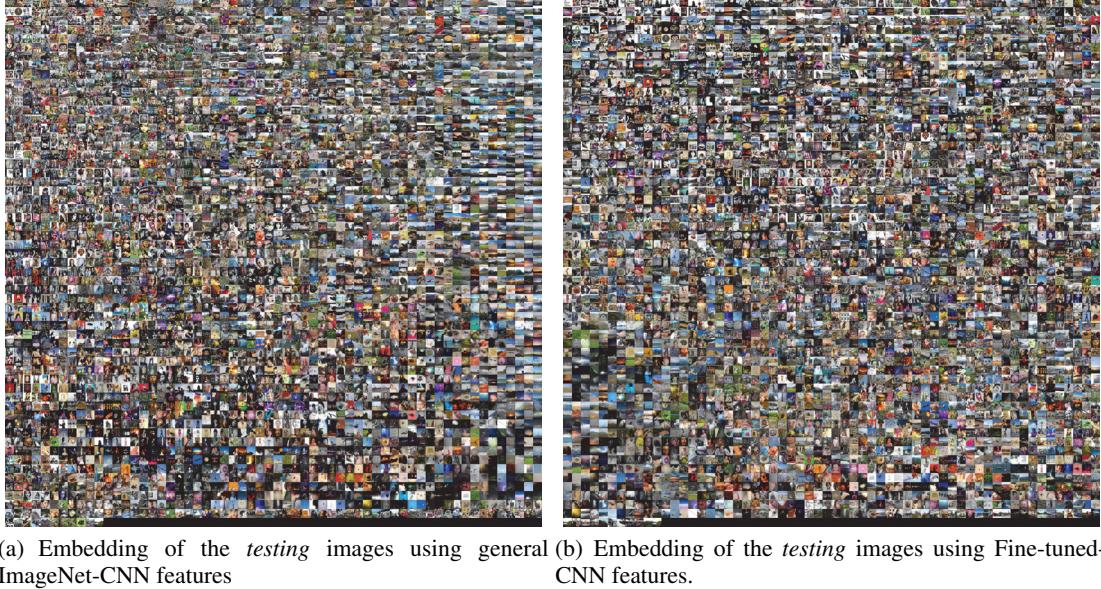


Figure 4: Visualization of learned filters for both ImageNet-CNN and the fine-tuned CNN (best viewed on screen with zoom).

Algorithms	Correct Samples	Accuracy
ImageNet-CNN	1120/3490	32.1%
Noisy-Fine-tuned-CNN	1600/3490	45.8%
Fine-tuned-CNN	2034/3490	58.3%

Table 3: Classification accuracy on the 15% randomly selected testing set labeled by the Amazon Mechanical Turk.

We compare the performance of deep features on visual emotion analysis with several other baseline features, including Wang et al. (Wei-ning, Ying-lin, and Sheng-ming 2006), Yanulevskaya et al. (Yanulevskaya et al. 2008), Machajdik and Hanbury (Machajdik and Hanbury 2010) and Zhao et al. (Zhou et al. 2014). Figures 5, 6 and 7 show the performance of these features on the three data sets respectively. Note that since emotion *anger* only contains 8 and 3 images in IAPS-Subset and Abstract Paintings data sets,

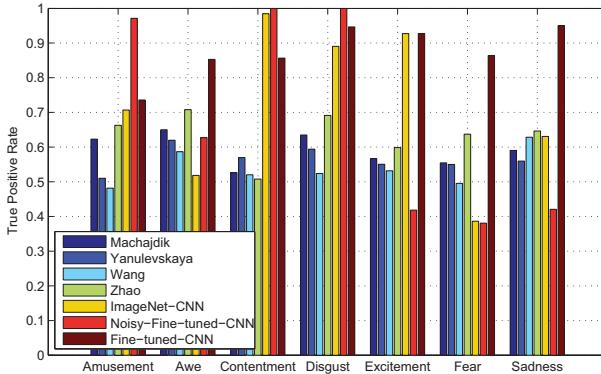


Figure 5: Per-class true positive rates of *Machajdik* (Machajdik and Hanbury 2010), *Yanulevskaya* (Yanulevskaya et al. 2008), *Wang* (Wei-ning, Ying-lin, and Sheng-ming 2006), *Zhao* (Zhou et al. 2014), *ImageNet-CNN*, *Noisy-Fine-tuned-CNN* and *Fine-tuned-CNN* on IAPS-Subset data set.

which are not enough to perform the 5-fold Cross Validation. We do not report the true positive rates for emotion *anger* on these two data sets.

It is interesting to find out that deep visual features significantly outperform the state-of-the-art manually crafted visual features in some emotion categories. However, the performance of using deep visual features are not consistent across the emotion categories at present. In particular, the performance of directly employing deep visual features from ImageNet-CNN and Noisy-Fine-tuned-CNN differ significantly among categories as well as across data sets. The performance of deep visual features from Fine-tuned-CNN is relatively more consistent. However, it has poor performance on emotions *Contentment* and *Fear* in the ArtPhoto data. These results suggest that it is still challenging to solve visual emotion analysis even with the state-of-the-art deep visual features. Meanwhile, the performance of deep visual features also suggests the promise of using CNNs in visual emotion analysis. Overall, this may encourage the development of more advanced deep architectures for visual emotion analysis, as well as development of other approaches.

## Conclusions

In this work, we introduce the challenging problem of visual emotion analysis. Due to the unavailability of a large scale well labeled data set, little research work has been published on studying the impact of Convolutional Neural Networks on visual emotion analysis. In this work, we are introducing such a data set and intend to release the data set to the research community to promote the research on visual emotion analysis with the deep learning and other learning frameworks. Meanwhile, we also evaluate the deep visual features extracted from differently trained neural network models. Our experimental results suggest that deep convolutional neural network features outperform the state-of-the-art hand-tuned features for visual emotion analysis. In addition, fine-tuned neural network on emotion related data sets can further improve the performance of deep neural net-

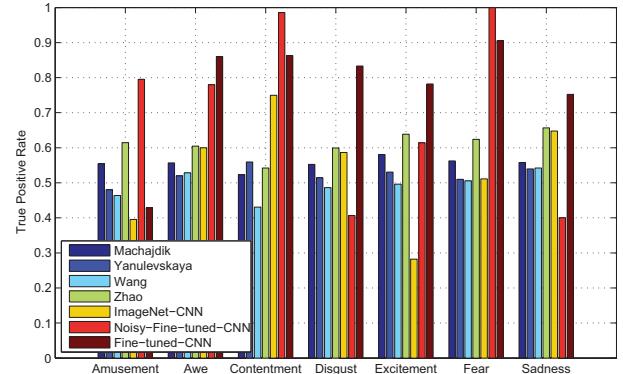


Figure 6: Per-class true positive rates of *Machajdik* (Machajdik and Hanbury 2010), *Yanulevskaya* (Yanulevskaya et al. 2008), *Wang* (Wei-ning, Ying-lin, and Sheng-ming 2006), *Zhao* (Zhou et al. 2014), *ImageNet-CNN*, *Noisy-Fine-tuned-CNN* and *Fine-tuned-CNN* on Abstract Paintings data set.

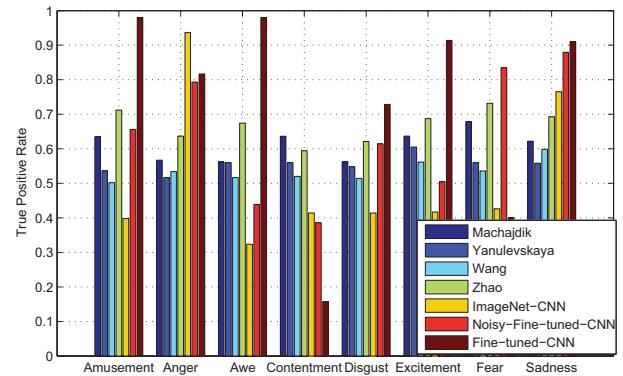


Figure 7: Per-class true positive rates of *Machajdik* (Machajdik and Hanbury 2010), *Yanulevskaya* (Yanulevskaya et al. 2008), *Wang* (Wei-ning, Ying-lin, and Sheng-ming 2006), *Zhao* (Zhou et al. 2014), *ImageNet-CNN*, *Noisy-Fine-tuned-CNN* and *Fine-tuned-CNN* on ArtPhoto data set.

work. Nevertheless, the results obtained in this work are only a start for the research on employing deep learning or other learning frameworks for visual emotion analysis. We will continue the collection of labeled data from AMT with a plan to submit additional 1 million images for labeling. We hope our visual emotion analysis results can encourage further research on online user generated multimedia content in the wild. Better understanding the relationship between emotion arousals and visual stimuli and further extending the understanding to valence are the primary future directions for visual emotion analysis.

## Acknowledgement

This work was generously supported in part by Adobe Research, and New York State CoE IDS. We thank the authors of (Zhao et al. 2014) for providing their algorithms for comparison.

## References

- Borth, D.; Ji, R.; Chen, T.; Breuel, T.; and Chang, S.-F. 2013. Large-scale visual sentiment ontology and detectors using adjective noun pairs. In *ACM MM*, 223–232. ACM.
- Cireşan, D. C.; Meier, U.; Masci, J.; Gambardella, L. M.; and Schmidhuber, J. 2011. Flexible, high performance convolutional neural networks for image classification. In *IJCAI*, 1237–1242.
- Deng, J.; Dong, W.; Socher, R.; Li, L.-J.; Li, K.; and Fei-Fei, L. 2009. Imagenet: A large-scale hierarchical image database. In *CVPR*, 248–255. IEEE.
- Grangier, D.; Bottou, L.; and Collobert, R. 2009. Deep convolutional networks for scene parsing. In *ICML 2009 Deep Learning Workshop*, volume 3.
- Hanjalic, A. 2006. Extracting moods from pictures and sounds: Towards truly personalized tv. *Signal processing magazine, IEEE* 23(2):90–100.
- Hinton, G. E.; Osindero, S.; and Teh, Y.-W. 2006. A fast learning algorithm for deep belief nets. *Neural computation* 18(7):1527–1554.
- Itten, J., and Van Haagen, E. 1973. *The art of color: the subjective experience and objective rationale of color*. Van Nostrand Reinhold New York, NY, USA.
- Jia, J.; Wu, S.; Wang, X.; Hu, P.; Cai, L.; and Tang, J. 2012. Can we understand van gogh’s mood?: learning to infer affects from images in social networks. In *ACM MM*, 857–860. ACM.
- Jia, Y.; Shelhamer, E.; Donahue, J.; Karayev, S.; Long, J.; Girshick, R.; Guadarrama, S.; and Darrell, T. 2014. Caffe: Convolutional architecture for fast feature embedding. *arXiv preprint arXiv:1408.5093*.
- Joshi, D.; Datta, R.; Fedorovskaya, E.; Luong, Q.-T.; Wang, J. Z.; Li, J.; and Luo, J. 2011. Aesthetics and emotions in images. *Signal Processing Magazine, IEEE* 28(5):94–115.
- Karayev, S.; Trentacoste, M.; Han, H.; Agarwala, A.; Darrell, T.; Hertzmann, A.; and Winnemoeller, H. 2013. Recognizing image style. *arXiv preprint arXiv:1311.3715*.
- Kavukcuoglu, K.; Sermanet, P.; Boureau, Y.-L.; Gregor, K.; Mathieu, M.; and LeCun, Y. 2010. Learning convolutional feature hierarchies for visual recognition. In *NIPS*, 5.
- Krizhevsky, A.; Sutskever, I.; and Hinton, G. E. 2012. Imagenet classification with deep convolutional neural networks. In *NIPS*, 1097–1105.
- Lang, P. J.; Bradley, M. M.; and Cuthbert, B. N. 1998. Emotion, motivation, and anxiety: brain mechanisms and psychophysiology. *Biological psychiatry* 44(12):1248–1263.
- Lang, P. J.; Bradley, M. M.; and Cuthbert, B. N. 1999. International affective picture system (iaps): Technical manual and affective ratings.
- Lang, P. J. 1979. A bio-informational theory of emotional imagery. *Psychophysiology* 16(6):495–512.
- LeCun, Y.; Boser, B.; Denker, J. S.; Henderson, D.; Howard, R. E.; Hubbard, W.; and Jackel, L. D. 1989. Backpropagation applied to handwritten zip code recognition. *Neural computation* 1(4):541–551.
- LeCun, Y.; Bottou, L.; Bengio, Y.; and Haffner, P. 1998. Gradient-based learning applied to document recognition. *Proceedings of the IEEE* 86(11):2278–2324.
- LeCun, Y.; Kavukcuoglu, K.; and Farabet, C. 2010. Convolutional networks and applications in vision. In *ISCAS*, 253–256. IEEE.
- Long, J.; Shelhamer, E.; and Darrell, T. 2014. Fully convolutional networks for semantic segmentation. *arXiv preprint arXiv:1411.4038*.
- Lu, X.; Suryanarayanan, P.; Adams, Jr., R. B.; Li, J.; Newman, M. G.; and Wang, J. Z. 2012. On shape and the computability of emotions. In *ACM MM*, 229–238. ACM.
- Lu, X.; Lin, Z.; Jin, H.; Yang, J.; and Wang, J. Z. 2014. Rapid: Rating pictorial aesthetics using deep learning. In *ACM MM*, 457–466. ACM.
- Machajdik, J., and Hanbury, A. 2010. Affective image classification using features inspired by psychology and art theory. In *ACM MM*, 83–92. ACM.
- Mikels, J. A.; Fredrickson, B. L.; Larkin, G. R.; Lindberg, C. M.; Maglio, S. J.; and Reuter-Lorenz, P. A. 2005. Emotional category data on images from the international affective picture system. *Behavior research methods* 37(4):626–630.
- Murray, N.; Marchesotti, L.; and Perronnin, F. 2012. Ava: A large-scale database for aesthetic visual analysis. In *CVPR*, 2408–2415. IEEE.
- Nicolaou, M. A.; Gunes, H.; and Pantic, M. 2011. A multi-layer hybrid framework for dimensional emotion classification. In *ACM MM*, 933–936. ACM.
- Valdez, P., and Mehrabian, A. 1994. Effects of color on emotions. *Journal of Experimental Psychology* 123(4):394.
- Van der Maaten, L., and Hinton, G. 2008. Visualizing data using t-sne. *Journal of Machine Learning Research* 9(2579–2605):85.
- Wei-ning, W.; Ying-lin, Y.; and Sheng-ming, J. 2006. Image retrieval by emotional semantics: A study of emotional space and feature extraction. In *Systems, Man and Cybernetics, IEEE International Conference on*, volume 4, 3534–3539. IEEE.
- Yanulevskaya, V.; Van Gemert, J.; Roth, K.; Herbold, A.-K.; Sebe, N.; and Geusebroek, J.-M. 2008. Emotional valence categorization using holistic image features. In *ICIP*, 101–104. IEEE.
- You, Q.; Luo, J.; Jin, H.; and Yang, J. 2015. Robust image sentiment analysis using progressively trained and domain transferred deep networks. In *AAAI*.
- Zhao, S.; Gao, Y.; Jiang, X.; Yao, H.; Chua, T.-S.; and Sun, X. 2014. Exploring principles-of-art features for image emotion recognition. In *Proceedings of the ACM International Conference on Multimedia*, MM ’14, 47–56. New York, NY, USA: ACM.
- Zhou, B.; Lapedriza, A.; Xiao, J.; Torralba, A.; and Oliva, A. 2014. Learning deep features for scene recognition using places database. In *Advances in Neural Information Processing Systems*, 487–495.